


Processing the World with H3 and PostGIS

Darafei Praliaskouski

PostGIS Project Steering Committee member
OpenStreetMap contributor

Founder at Maumap.com Geospatial consulting
Previously at Kontur.io

OpenAerialMap / OpenDroneMap fan
Lately Meshtastic LoRa radio planning
Living in Batumi, Georgia





Georgia has mountains



Mountains are fun for radio

Telecom towers are always on top.

You can put yours nearby.




Meshtastic

An open source, off-grid, decentralized, mesh network built to run on affordable, low-power devices.

<https://meshtastic.org/>

Runs on LoRa,
which you can meet in IoT
Devices as low as \$12...
Why not put them
everywhere?



What we learned deploying our LoRa network

- Surprisingly long range on direct visibility - we have 56km links.
- Still works in urban setting. More people, more noise, more volunteers, more repeaters.
- It's easy to put a solar powered node on the pole on reasonable height.
- Largest issue is logistics of getting node there and picking a spot where it's useful.
- There are people who climb mountains for fun and don't mind putting your node there.
- Having free direct messaging with other engineers is fun :)

So, today we'll look at the question - **where** do we put our nodes?


H3

Hierarchical Geographical Index

Mostly hexagonal

Useful for statistical analysis

Optimized for analyst's mental health



In the H3 indexing system,
how many edges does a hexagon have on average?

At resolution 1, average hexagon has **6.41** edges

At the odd H3 resolutions
some hexagons get extra **distortion vertices**
when they belong to multiple icosahedral faces

In addition to **560** hexagons,
resolution 1 have:
• **240** heptagons
• **30** octagons
• **12** decagons


KONTUR

KONTUR

Pentagons in H3 hexagonal grid.
All of them are
centered on water +




#30DayMapChallenge




Adventures of Polygons traveling between different Earth models

Flat sheet to Globe : points you may expect inside are now outside when the segments follow great circle rather than rhumb lines.



Globe to Flat sheet to Cylinder :
the polygon is now inside out on every rotation of the cylinder.

Flat sheet to Cylinder : edges of a flat map can sometimes be seen as a boundary inside the polygon



h3_pg

Best set of extensions for supporting h3 in Postgres.

<https://github.com/zachasme/h3-pg>

Created and maintained by Zacharias Knudsen

Thanks a lot for accepting our pull requests :)



Start by choosing your uncertainty

Most of my projects are in one of two resolutions:

- Resolution 8:
 - “walking distance”
 - ~400m edge length
 - the planet will likely fit into your laptop
- Resolution 11:
 - “specific building”
 - ~ 30m edge length
 - can use H3 for pathfinding

$$\sigma_x \sigma_p \geq \frac{\hbar}{2}$$

If you think you need more detail, likely H3 isn't the best tool for your task.

Formalizing task

I'd love to **find locations** for growing a network of radio repeaters in **Georgia**.

I want every next one to **grow** the network towards new **populated area**.

Resolution 8 should be okay for first pass.

Let's assume that 60 km links in **line of sight** are working every time.

DEM will be enough - we're not looking closely at cities this time.

I want to get to the installation site **by car**.

Model we build shouldn't contradict existing **ground truth** cases we observe.




-Ground truth data

Batumi can see Poti

Poti sees Gomis Mta

Batumi does not see Gomis Mta

Neither Poti nor Batumi see Tbilisi.



What data will we need?

Some existing nodes

Elevation raster

Roads

Country boundary

Population

What data will we need?

Some existing nodes points

Elevation raster raster

Roads lines

Country boundary polygon

Population h3 grid

What data will we need?

Some existing nodes points sketch in geojson.io

Elevation raster raster GEBCO 400m global grid

Roads lines OpenStreetMap

Country boundary polygon OpenStreetMap

Population h3 grid Kontur Population

GeoJSON points to H3

```
$ ogr2ogr -f PostgreSQL PG:@"" data/in/existing_mesh_nodes.geojson -nln  
mesh_initial_nodes -nlt POINT -lco GEOMETRY_NAME=geom -overwrite -a_srs  
EPSG:4326
```

```
-- Create seed towers table projected to H3 resolution 8  
  
create table mesh_initial_nodes_h3_r8 as  
  
select  
    h3_latlng_to_cell(geom, 8) as h3,  
    string_agg(coalesce(name, 'seed'), ', ') as name,  
    ST_Collect(geom) as geom  
from mesh_initial_nodes  
group by 1;
```

OpenStreetMap osmium import one-liner


```
osmium export -i sparse_mem_array -c osmium.config.json -f
pg data/in/osm/georgia-latest.osm.pbf -v --progress
| psql -1 -c "create table osm_georgia(geog geography,
osm_type text, osm_id bigint, version int, osm_user text, ts
timestamptz, way_nodes bigint[], tags jsonb); alter table
osm_georgia alter geog set storage external, alter osm_type
set storage main, alter osm_user set storage main, alter
way_nodes set storage external, alter tags set storage
external, set (fillfactor=100); copy osm_georgia from stdin
freeze;"
```

OpenStreetMap: Boundaries of Georgia

```
create table georgia_boundary as
select ST_Union(geog::geometry) as geom
from osm_georgia
where tags ? 'boundary'
and tags -> 'boundary' = 'administrative'
and tags -> 'admin_level' = '2'
and (
    coalesce(tags -> 'name', "") ilike 'georgia'
    or coalesce(tags -> 'name:en', "") ilike 'georgia'
);

```

```
create table georgia_convex_hull as
select ST_ConvexHull(ST_Collect(geom)) as geom
from georgia_boundary;
```



GEBCO: Global Bathymetric Chart of the Oceans

British Oceanographic Data Centre

very kindly provides elevation dataset that has both depth of the oceans and elevation above sea level.

15 arc-second interval grid.

$111.3 \text{ km} * 0.0041667 \approx 463.8 \text{ m.}$

It nicely matches H3 Resolution 8.



Importing GEBCO

```
$ raster2pgsql -I -C -M -s 4326 -t auto data/mid/gebco/*.tif  
gebco_elevation_all | psql -q  
  
$ psql -v table_name=gebco_elevation \  
      -v table_name_h3=gebco_elevation_h3_r8 \  
      -v item_name=ele \  
      -v aggr_func=avg \  
      -v resolution=8 \  
      -v clip_table=georgia_convex_hull \  
      -f scripts/raster_values_into_h3.sql
```

raster_values_into_h3.sql

```
-- Converts raster pixels to H3 aggregates using provided parameters
-- Expect psql variables: table_name, table_name_h3, item_name, aggr_func, resolution, clip_table
drop table if exists :table_name_h3;
create table :table_name_h3 as
with clip as (
    select geom from :clip_table
)
select
    H3,
    :resolution::int as resolution,
    agg_val as :item_name
from (
    select
        h3_latlng_to_cell(p.geom::geography, :resolution) as h3,
        :aggr_func(p.val) as agg_val
    from :table_name t,
        Clip,
        LATERAL ST_PixelAsCentroids(ST_Clip(t.rast, clip.geom, true)) as p
    group by 1
) s;
```

GEBCO in H3 - 10 color classes


Roads for motor vehicles

```
create table georgia_roads_geom as
select  geog::geometry as geom,
        tags ->> 'highway' as highway
from osm_georgia
where tags ? 'highway'
    and (tags ->> 'highway') = any (
        Array[
            'motorway', 'motorway_link', 'trunk', 'trunk_link',
            'primary', 'primary_link', 'secondary', 'secondary_link', 'tertiary', 'tertiary_link',
            'unclassified', 'residential', 'living_street', 'service', 'road'
        ]
)
and coalesce(
    lower(tags ->> 'motor_vehicle'),
    lower(tags ->> 'motorcar'),
    lower(tags ->> 'vehicle'),
    lower(tags ->> 'access'),
    'yes'
) not in ('no', 'private')
and ST_GeometryType(geog::geometry) in ('ST_LineString', 'ST_MultiLineString');
```

Roads to Segments to Points to H3

```
create table roads_h3_r8 as
select
    h3_latlng_to_cell(ST_StartPoint(seg_geom)::geography, 8) as h3,
    sum(ST_Length(seg_geom)::geography)) as road_length_m
from (
    select (ST_DumpSegments(
        ST_Segmentize(
            geom::geography,
            200)::geometry)
        ).geom as seg_geom
    from georgia_roads_geom
) segments
where seg_geom is not null
group by 1;
```

Roads as binary H3 cells



Population data in H3

```
$ ogr2ogr -f PostgreSQL PG:""  
data/mid/population/kontur_population_20231101.gpkg -nln  
kontur_population -nlt MULTIPOLYGON -lco GEOMETRY_NAME=geom  
-overwrite -t_srs EPSG:4326
```

It's already in H3.

Combine everything!


```
create table mesh_surface_h3_r8 (
    h3                                h3index primary key,
    geom geometry generated always as (h3_cell_to_boundary_geometry(h3))
stored,
    centroid_geog public.geography generated always as
(h3_cell_to_geometry(h3)::public.geography) stored,
    ele double                          precision,
    has_road                           boolean default false,
    population                         numeric,
    has_tower                           boolean default false,
    has_reception                      boolean,
    can_place_tower                    boolean,
    visible_uncovered_population      numeric,
    distance_to_closest_tower         double precision
);
```

Line of Sight Visibility function

```
create or replace function h3_los_between_cells(h3_a h3index,
h3_b h3index)
    returns boolean
language plpgsql
cost 100000
as
$$
...
/*
 * 1) h3_grid_path_cells(norm_src, norm_dst) gives us
the H3 cells along
    *   the path between the towers (inclusive).
 * 2) We join that to mesh_surface_h3_r8 to get
elevations.
    * 3) For each intermediate cell:
        *   frac = position of cell centroid along the line
[0..1]
        *   los_height(frac) = ele_src + (ele_dst - ele_src)
* frac
        * If terrain elevation >= los_height at any point, LOS
is blocked.
    */
with path as (
    select h3_grid_path_cells(norm_src, norm_dst) as h3
),
```

```
samples as (
    select
        c.h3,
        ST_LineLocatePoint(
            line_geom,
            ST_PointOnSurface(c.geom)
        )::double precision as frac,
        c.ele
    from path p
    join mesh_surface_h3_r8 c
        on c.h3 = p.h3
    where p.h3 not in (norm_src, norm_dst)
)
select exists (
    select 1
    from samples s
    where s.frac > 0
        and s.frac < 1
        and (
            ele_src + (ele_dst - ele_src) * s.frac
        ) <= coalesce(s.ele, 0)
)
into has_blocker;
is_visible := not coalesce(has_blocker, false);
return is_visible;
```

Initial coverage: testing links



Greedy loop to pick growth direction

1. Refresh reception cache

For each surface cell `s`:

`has_reception(s) = exists tower t within 60 km with LOS`

2. Enable tower candidates

`can_place_tower = true where`

`has_road and not has_tower and distance in [5 km, 60 km].`

3. Score candidates by uncovered population

For each candidate `c`

`visible_uncovered_population(c) = sum population(t)`

over cells `t` that are

`uncovered + within 60 km + LOS(c, t).`


4. Greedy pick + update

Choose `c*` with max `visible_uncovered_population`.

Insert tower at `c*`, store iteration.

Around `c*` (60 km): update `distance_to_closest_tower`, kill candidates closer than 5 km, invalidate caches.

Result - ideal visibility-based LoRa mesh



Useful links

<https://github.com/Komzpa/h3-mesh-placement>

<https://github.com/konturio/geocint-runner/blob/main/DOCUMENTATION.md>

<https://meshtastic.org/>

<https://github.com/zachasme/h3-pq>

<https://www.gebco.net/>

https://en.wikipedia.org/wiki/Nyquist%E2%80%93Shannon_sampling_theorem

Let's be in touch

Darafei Praliaskouski

darafei@maumap.com